

Jason (Guoxuan) Xu

No Sponsorship Required | (+1) 909-760-9522 | guoxuan.xu8@gmail.com | [Linkedin](#) | [Github](#)

EDUCATION

University of California San Diego

BS. Data Science (*GPA: 4.00, Provost Honors*)

La Jolla, CA

Sep 2022 – Jun 2026 (Expected)

INDUSTRY EXPERIENCE

Cadre AI (AI Consulting) - *AI Engineering Intern*

July 2025 – Aug 2025

- Developed a **cloud-native, serverless AI agent framework** leveraging **Azure Functions, Python, and Blob Storage**, enabling traditional firms to automate document handling and storage without maintaining infrastructure.
- Designed and deployed **5 production-ready AI agents** as modular pipelines, cutting document processing time from **2 hours to 10 minutes** for 100+ employees while ensuring **0% error rate** and compliance with formatting standards.
- Engineered reusable pipelines with **HTTP triggers & serverless orchestration**, reducing AI agent development and deployment cycles from weeks to hours across **4 enterprise clients**.
- Led **end-to-end system design and implementation**, collaborating with the AI manager to align requirements, enforce compliance, and ensure seamless adoption across client workflows.

LLM Strategies (AI startup) - *Data Engineering Intern*

Dec 2024 – Jun 2025

- Designed and implemented distributed pipelines with **Celery + Docker**, enabling scalable ingestion of millions of patent and trademark records from **REST APIs, HTML scrapers, and XML feeds**.
- Built production-ready services for high-accuracy document image extraction (USPTO API + OCR fallback + Gemini API fallback), achieving **99.2% accuracy** across 500+ multi-format court documents.
- Optimized system performance by implementing streaming XML parsing (*lxml.iterparse*) for weekly 10k+ records and reducing redundant metadata by **85%** (2.5M → 362k) through algorithmic cleaning and normalization.
- Developed a full-stack data platform by integrating ETL pipelines with a **PostgreSQL → Data Warehouse backend**, exposing analytics through **Flask APIs**, and building **React dashboards** for real-time trend visualization.

The Wu lab, Scripps Research - *Software Engineering Intern*

Nov 2024 – Feb 2025

- Developed a **FastAPI service** to expose genomic datasets as RESTful APIs, integrating **PostgreSQL** for structured metadata and **S3 storage** for raw files, enabling global researchers to query standardized biological data.
- Engineered high-performance **Python parsers** with **Ixml** and generator-based pipelines to normalize raw genomic datasets, applying rule-based validation to ensure data quality while preventing memory bottlenecks on large files.

PROJECT EXPERIENCE

Trading Capture System – Execution Module - *Group Project*

Jul 2025 – Present

- Developed the **execution module** of a distributed trading system with a **microservices backend** and **React frontend**, enabling secure execution management and real-time communication.
- Implemented **RESTful APIs** to process **FIX messages**, ensuring reliable event-driven communication across distributed services.
- Designed and optimized **MySQL schemas**, and integrated **Redis caching** for token management (logout & refresh), improving query performance and authentication efficiency.
- Built frontend features with **React + Redux** to create, update, and delete execution records, providing real-time execution tracking and improved user experience.
- Implemented execution validation logic by enforcing field constraints and dependency checks, ensuring data integrity.
- Engineered **Kafka-based asynchronous reporting**, enhancing system responsiveness and reliability in a distributed environment.

Climate Emulation with Deep Learning - *Kaggle Competition*

May 2025 – Jun 2025

- Engineered a **U-Net model with CoordConv2d and Residual Layers** in PyTorch, achieving **Top 10 (9/85)** in forecasting surface temperature and precipitation.

RESEARCH EXPERIENCE

Student Researcher - *Data Science & Learning Lab @ UCSD*

Mar 2024 – Jul 2024

- Enhanced the [Pandas Tutor visualization tool](#) by improving DataFrame introspection logic, enabling selective rendering of large datasets to reduce memory usage and speed up interactivity in-browser.
- Published as **first co-author** at **ACM SIGCSE**, analyzing data analysis behaviors while contributing to the design and optimization of large-scale visualization features.

SKILLS & TOOLS

Languages & Frameworks: Java, Python, SQL/PostgreSQL, JavaScript, R, Shell, PyTorch, scikit-learn, React

Tools & Platforms: Git, Docker, Kubernetes, Redis, MySQL, Kafka, RESTful APIs, Azure, Flask, FastAPI