

# Guoxuan Xu

No Sponsorship Required | (+1) 909-760-9522 | guoxuan.xu8@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## EDUCATION

**University of California San Diego** | B.S. Data Science (*GPA: 3.97, Provost Honors*) Sep 2022 – Jun 2026 (Expected)  
**Stanford University** | M.S. Computer Science Sep 2026 – Mar 2028 (Expected)

## INDUSTRY EXPERIENCE

**Samsung - Incoming Software Development Engineer Intern** Summer 2026  
**Cadre AI (AI Consulting) - Part-Time AI Engineer** July 2025 – Feb 2026

**Project:** Serverless AI Agent Automation Platform

- Developed a **cloud-native, serverless AI agent framework** using **Azure Functions**, **Python**, and **Blob Storage**, enabling document automation with zero infrastructure overhead.
- Designed and deployed **5 modular AI agents** (Sales Proposal, Preflight, SOP, Compliance, Q&A), reducing doc processing time from **2h** → **10m** for **100+** employees with **0%** formatting/compliance errors.

**Project:** MCP Proxy for Enterprise API Integrations

- Built a **FastAPI-based MCP proxy** enabling **ChatGPT** to securely access **BigQuery** and **Zoom** through **OAuth 2.0**, improving dataset exploration and workflow automation for **20+ internal users**.
- Implemented proactive token refresh and request forwarding for long-running ChatGPT sessions, reducing manual BigQuery console navigation by **~70%**.

**Project:** RAG Admin Portal & User Management System

- Designed and implemented a **full-stack RAG admin portal** with token-validated signup, RBAC, role management, account deactivation, secured authentication endpoints, and **frontend caching** to improve access control and page responsiveness.

**Project:** Intelligent Email Automation Platform (Azure)

- Architected an **event-driven email automation system** using **FastAPI** and **Azure App Service**, integrating Microsoft Inbox webhooks, OCR, and LLM pipelines to extract structured data and automate validation and response drafting.

**LLM Strategies (AI startup) - Data Engineering Intern** Dec 2024 – Jun 2025

- Built distributed ingestion services and Flask microservices with **Celery**, **Redis**, and **Docker** to process **10M+ patent and trademark records** from APIs, scrapers, and XML feeds.
- Developed OCR, Gemini, SigLIP, and **Qdrant-based extraction** and **vector search pipelines**, achieving **99.2% accuracy** on **500+ legal documents** and reducing redundant metadata by **85%**.
- Improved pipeline reliability through **asynchronous task execution**, **retry handling**, and **containerized deployment workflows**.

## PROJECT & RESEARCH EXPERIENCE

**Trading Capture System – Execution Module** | *Spring Boot, Kafka, FIX, MySQL, Redis, React* Jul 2025 – Present

- Developed core modules of a **distributed trading system** using microservices (Java/Spring Boot) and React, enabling secure execution management and real-time communication across services.
- Implemented **FIX-driven REST APIs** and **Kafka-based asynchronous reporting**, ensuring reliable event-driven message flow and improving system responsiveness in distributed environments.
- Designed and optimized **MySQL** schemas and integrated **Redis** caching for token/session management, improving query latency and authentication efficiency.
- Built **React + Redux** interfaces for creating, updating, and tracking execution records, and implemented validation logic with field-level constraints to ensure data integrity and smooth user experience.

**ICU Sepsis Risk Prediction Platform** | *Django, AWS RDS, Random Forest, HTML/CSS* Jul 2025 – Present

- Built a **Django MVT web app** for ICU patient monitoring and sepsis risk prediction using historical patient data and Random Forest inference.
- Integrated **AWS RDS** with ORM-based data access and cached prediction outputs to reduce repeated inference latency.
- Optimized model artifact size for cloud deployment and developed interactive pages for patient status tracking and risk score review.

**AI-Powered Sprint Review Assistant** | *Claude API, Prompt Engineering, Full-Stack Development* Mar 2024 – Jul 2024

- Built a full-stack productivity app integrating **Claude API** and category-specific prompt templates to convert weekly notes into structured sprint reviews, including achievements, blockers, reflections, action plans, and historical tracking.

**Student Researcher - Data Science & Learning Lab @ UCSD** Mar 2024 – Jul 2024

- Enhanced the [Pandas Tutor visualization tool](#) by improving DataFrame introspection for selective rendering of large datasets. Published this work as **first co-author** at ACM SIGCSE (paper: <https://dl.acm.org/doi/10.1145/3641554.3701933>).

## SKILLS & TOOLS

**Languages & Frameworks:** Java, Spring Boot, Python, SQL/PostgreSQL, JavaScript, React, Flask, FastAPI, Django, R, Shell

**Tools & Platforms:** Git, Docker, Kubernetes, Redis, MySQL, Kafka, RESTful APIs, Azure, AWS RDS, Qdrant, Celery